

NEURAL NETWORK OPTIMIZATION USING PARALLEL COMPUTING

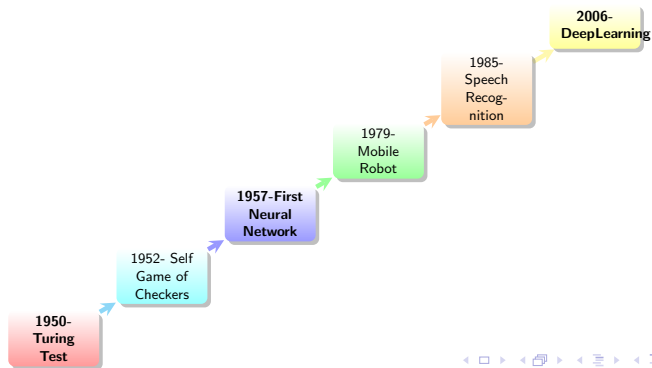
PROJECT FOR MASTER'S DEGREE IN HPC

Saumya Bhatnagar


Trinity College Dublin

October 2, 2018

■ Fourth Industrial Revolution & Evolution of Machine Learning

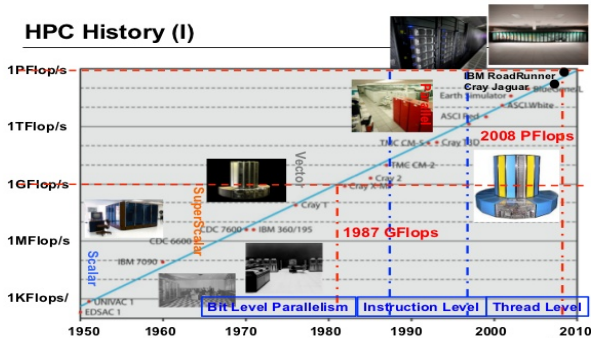


- Fourth Industrial Revolution & Evolution of Machine Learning
- End of Moore's law and Quantum Mechanics(1)



- Fourth Industrial Revolution & Evolution of Machine Learning
- End of Moore's law and Quantum Mechanics(1)
- Evolution of HPC (2)

HPC History (I)



- Fourth Industrial Revolution & Evolution of Machine Learning
- End of Moore's law and Quantum Mechanics(1)
- Evolution of HPC (2)
- Trends of HPC (2)

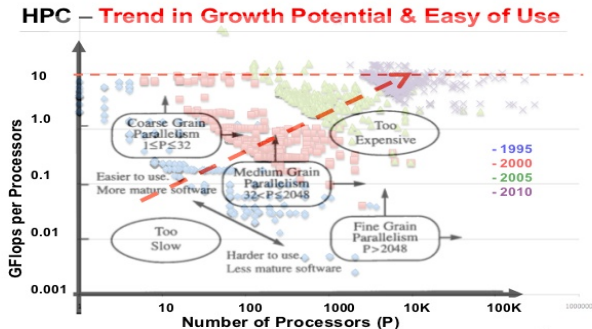


TABLE OF CONTENTS

1 INTRODUCTION & BACKGROUND


2 SOFTWARE DEVELOPMENT

- Optimization Method
- Parallel Computation
- Implementation
- Optimization Method - Quantum Inspired

3 ANALYSIS & CONCLUSIONS

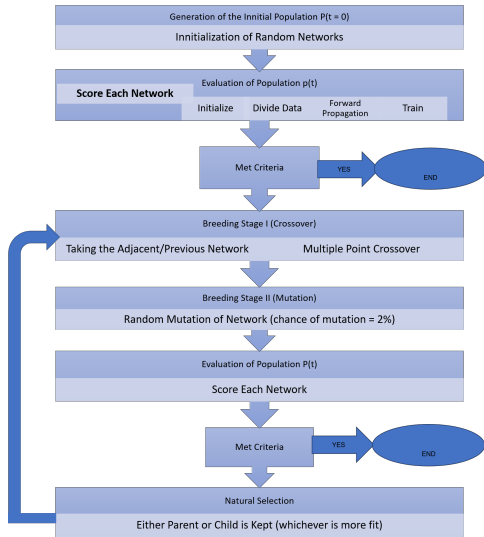
- Sequential Algorithm Analysis
- Sequential Algorithm vs Parallel Algorithm
- Parallel Algorithm - Island Model
- Efficiency & Speedups
- Quantum Inspired Algorithm

- 1 Compared accuracy of Neural Nets
- 2 Single Population Fine Grained Algorithms
- 3 Multiple-population Coarse Grained Algorithms(3)
- 4 Sequential Time Consumption
- 5 Mpi4py, Keras, Tensorflow
- 6 Analyzed TCHPC Clusters
- 7 Memory Leak



- To access best set of hyper-parameters
- Convolutional Neural Network
- Image Processing
- Machine Learning Benchmark : CIFAR10
- Initiate with randomly selected combination of hyper-parameters
- Sample space kept same with every selection
- MPI parallelism
- Scheduling, Re-Scheduling by resuming the code based on Node availability using Slurm Workload Manager
- Keras has built-in support for multi-GPU data parallelism
- Boyle Cluster: ldd (GNU libc) 2.17

Optimization Method




- Inspired by GA (4)
- To crossbreed networks
- Relative accuracy of hyper-parameters set is analyzed

Dataset	Data-points	Sequential Time (per generation)
CIFAR10	10k	15 minutes
80M Tiny Images	80M	20k minutes

Number of Processors $< 32 \Rightarrow$ Implemented Course Grained Parallelism
(6)

Two Parallel Implementations of GA

- Single Population Fine Grained Algorithms** : using 4, 6, 8, 10, 12 processors




Dataset	Data-points	Sequential Time (per generation)
CIFAR10	10k	15 minutes
80M Tiny Images	80M	20k minutes

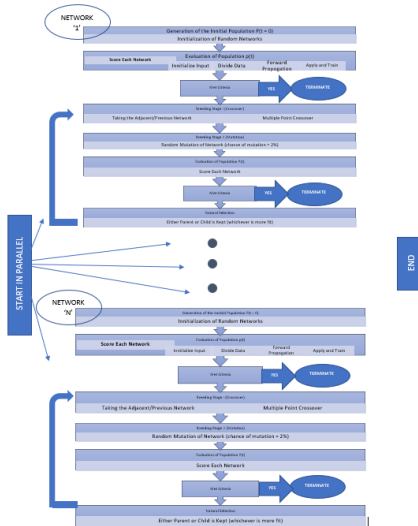
Number of Processors $< 32 \Rightarrow$ Implemented Coarse Grained Parallelism (6)

Two Parallel Implementations of GA

- 1 **Single Population Fine Grained Algorithms** : using 4, 6, 8, 10, 12 processors
- 2 **Multiple-population Coarse Grained Algorithms** : with inter-island (every 5 generations) and intra-island breeding
 - population size of 4 in each island, using 12 processors



Parallel Computation



Single Population Fine Grained Algorithms

Send/Recv Model

One Network Per Node

Inter-Network Crossover

Multiple-population Coarse Grained Algorithms

MPI Communicators

One Communicator Per Island

Inter-Island & Intra-Island Crossover

MPI Initialization Testing

```
MPI.Init()
print(MPI.Is_initialized())
print(MPI.Is_finalized())
```

Non-Blocking Exchange

```
def nonBlockingExchange(self, data):
    reqSend1 = self.comm.isend(data,
                                dest=((self.size+self.rank+1)\%self.size),
                                tag=self.rank)
    reqRecv2 = self.comm.irecv(source=
                                ((self.size+self.rank-1)\%self.size),
                                tag=self.rank-1)
    dataPrev = reqRecv2.wait()
    reqSend1.wait()
    return dataPrev
```

MPI Gather

```
recvdata = self.comm.Gather(data, root = 0)
```

MPI Broadcast

```
data = self.comm.bcast(data, root = 0)
```

Get Island

```
self.subGroup = self.rank // groupSize
self.subComm = MPI.Comm.Split(
    self.comm, self.subGroup, self.rank)
```

Get Island Details

```
self.subSize, self.subRank =
    self.subComm.Get_size(), self.subComm.Get_rank()
```

Qubit in super-position state, increases sample space(5)

Hadamard gate

```
r2=math.sqrt(2.0)
```

```
h=np.array([[1/r2, 1/r2],[1/r2,-1/r2]])
```

Rotation Q-gate

```
rot1=float(math.cos(theta));rot2=-float(math.sin(theta))
```

```
rot3=float(math.sin(theta));rot4=float(math.cos(theta));
```

Quantum Population Vector

```
# alpha squared
```

```
self.qpv[i,j,0]=np.around(2*pow(AlphaBeta[0],2),2)
```

```
# beta squared
```

```
self.qpv[i,j,1]=1-self.qpv[i,j,0]
```

Quantum Measure

- Search space consists of 5 hyper-parameters: activation functions, optimizers, hidden layers, nodes and dropout (7)
- Six activation functions included are: sigmoid, elu, selu, relu, tanh, hard_sigmoid
- Six Optimizers include sgd, adagrad, adadelta, adam, adamax, nadam.
- Hidden layers range from 1-15.
- Nodes/neurons range from 4-128.
- Dropouts range from 0.1 to 0.5.


SAMPLES TESTED

Activation	Optimizer					
	'adadelata'	'adagrad'	'adam'	'adamax'	'nadam'	'sgd'
'elu'	■					■
'hard_sigmoid'			■	■	■	■
'relu'	■	■	■	■		■
'selu'		■	■		■	■
'sigmoid'	■	■	■	■		■
'tanh'	■		■	■	■	■

- Activation functions: hard_sigmoid and elu couldn't come in any set of combinations with the optimizer - adagrad
- Uneven Testing: some combinations have been tested more frequently

SAMPLES TESTED





Activation	Optimizer				
	'adadelta'	'adagrad'	'adam'	'adamax'	'sgd'
'elu'	■				■
'hard_sigmoid'	■		■	■	■
'relu'	■	■	■	■	■
'selu'	■	■	■	■	■
'sigmoid'	■	■	■	■	■
'tanh'	■		■	■	■



- Activation functions: `hard_sigmoid` and `elu` couldn't come in any set of combinations with the optimizer - `adagrad`
- Uneven Testing: some combinations have been tested more frequently
- Increasing the number of networks, increased the possible combinations
- Higher neurons and lower layers factored in high accuracies

Sequential Algorithm vs Parallel Algorithm

FITNESS COMPARISON




The convergence of parallel implementations is better than that of sequential implementation.

Sequential Algorithm vs Parallel Algorithm


Fitness Comparison - 8 Networks (Serial)




Fitness Comparison - 8 Networks (Parallel)




Fitness Comparison - 10 Networks (Serial)




Fitness Comparison - 10 Networks (Parallel)



Fitness Comparison - 12 Networks (Serial)







Fitness Comparison - 12 Networks (Parallel)









Sequential Algorithm vs Parallel Algorithm

TIME COMPARISON




- The time taken by sequential code increases with generations.
- Time taken by Sequential code is in hours while the same can be done in minutes using MPI parallelism

Sequential Algorithm vs Parallel Algorithm





MULTIPLE-POPULATION COARSE GRAINED ALGORITHMS

Time Lapse Comparison -4 Networks (Parallel vs Island Parallel vs Sequential)



The Island model showed speed-up


SUPER-LINEAR BEHAVIOR



SUPER-LINEAR BEHAVIOR & MEMORY LEAK



Super-linear speedup




MEMORY LEAK

types	Running for 1 generation		Running for 5 generations		Times increment in number of objects
	# objects	total size	# objects	total size	
<class 'numpy.ndarray'	3	117.23 MB	3	117.23 MB	NA
<class 'tuple'	62126	5.69 MB	482777	44.19 MB	7.8
<class 'list'	44141	4.30 MB	250113	24.86 MB	5.7
<class 'dict'	10997	2.52 MB	87719	20.02 MB	8.0
<class 'int'	83692	2.23 MB	630694	16.84 MB	7.5
<class 'tensorflow.core.framework.node_def_pb2.NodeDef'			28057	2.35 MB	Internally saved from first generation in keras, and then incremented with generations
<class 'tensorflow.python.framework.tensor_shape.Dimension'			29688	1.59 MB	
<class 'tensorflow.python.framework.ops.Tensor'			29401	1.57 MB	
<class 'tensorflow.python.framework.tensor_shape.TensorShape'			29361	1.57 MB	
<class 'tensorflow.python.framework.ops.Operation'			28057	1.50 MB	
<class 'str'	14656	1.05 MB	17258	1.24 MB	1.18


Approx 100MB memory leak in 5 generations

Quantum Inspired Algorithm


Time taken by Parallel Implementation vs Quantum Inspired Algorithm



Fitness Achieved by Parallel Implementation vs Quantum Inspired Algorithm



Quantum Inspired Algorithm



The trend of sum of Quantum Inspired for Generation. Color shows details about Data.

REFERENCES I

- [1] Thomas N Theis and H-S Philip Wong. The end of moore's law: A new beginning for information technology. *Computing in Science & Engineering*, 19(2):41–50, 2017.
- [2] John L Hennessy and David A Patterson. *Computer architecture: a quantitative approach*. Elsevier, 2011.
- [3] Erick Cantú-Paz. A survey of parallel genetic algorithms. *Calculateurs paralleles, reseaux et systems repartis*, 10(2):141–171, 1998.
- [4] John J Hopfield and David W Tank. “neural” computation of decisions in optimization problems. *Biological cybernetics*, 52(3): 141–152, 1985.

REFERENCES II

- [5] Ling Wang, Fang Tang, and Hao Wu. Hybrid genetic algorithm based on quantum computing for numerical optimization and parameter estimation. *Applied Mathematics and Computation*, 171(2): 1141–1156, 2005.
- [6] Jian-Ming Li, Xiao-Jing Wang, Rong-Sheng He, and Zhong-Xian Chi. An efficient fine-grained parallel genetic algorithm based on gpu-accelerated. In *Network and parallel computing workshops, 2007. NPC workshops. IFIP international conference on*, pages 855–862. IEEE, 2007.
- [7] Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT press Cambridge, 2016.